

Introduction to Stan

Wei Wang
(mostly borrowed from Ben Goodrich though)

G6103 BDA: October 1, 2014

Outline

- 1 Introduction
- 2 Installation
- 3 Bayesian Inference
- 4 Stan
- 5 Examples

Plan for Today

- 1 Install some flavor of Stan on everyone's laptop
- 2 Briefly talk about Bayesian inference and MCMC
- 3 Overview of the Stan modeling language
- 4 Examples of using Stan

Links

- <http://mc-stan.org/> has everything
- Google Groups:
 - low-volume release announcements:
<https://groups.google.com/forum/?fromgroups#!forum/stan-announce>
 - for help with your models / configuration problems:
<https://groups.google.com/forum/?fromgroups#!forum/stan-users>
 - if you are interested in contributing to Stan:
<https://groups.google.com/forum/?fromgroups#!forum/stan-dev>

Flavors of Stan

- “Stan” is a catch-all term that includes
 - `libstan`, a library for statistics and optimization
 - `stanc`, a parser for the Stan language (convert Stan code to C++ code)
 - interfaces to `libstan` and `stanc`
 - `CmdStan`, for use via a command-line shell
 - `RStan`, for use via the R language
 - `PyStan`, for use via the Python language
 - `MStan` (MATLAB), `StataStan` (Stata), etc. are in progress
 - the Stan community
- Install the interface that is most comfortable for you (by which I meant `rstan`...)

PyStan

- Requires Python 2.7+; see <https://pystan.readthedocs.org/>
- Helps to have matplotlib
- Windows: <https://pystan.readthedocs.org/en/latest/windows.html>
- Linux or OS X **prior to Mavericks**: Use pip via shell

```
sudo pip install pystan
```

- OS X Mavericks: Install from source via Terminal
 - Requires Cython and NumPy

```
wget https://github.com/stan-dev/pystan/archive/2.2.0.1.zip
unzip 2.2.0.1.zip
cd pystan-2.2.0.1
sudo export MAKEFLAGS = "-j4" # or another number besides 4
sudo \
ARCHFLAGS=-Wno-error=unused-command-line-argument-hard-error-in-future \
sudo python setup.py install
export \
ARCHFLAGS=-Wno-error=unused-command-line-argument-hard-error-in-future
cd ..
```

RStan

- Not on CRAN yet
- Somewhat involved installation procedure
- <https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started>
- Can utilize Amazon EC2 http://www.louisaslett.com/RStudio_AMI/ but “micro” instances have too little RAM

CmdStan

- Possible to access Stan from command-line
- Probably mainly of interest to potential developers

```
git clone https://github.com/stan-dev/cmdstan  
make /path/to/stanfile-without.stan-extension
```


Bayes' Theorem

- Let $\vec{\theta}$ be a vector of unknown parameters
- Let $f(\cdot)$ be a PDF (continuous) or PMF (discrete)
 - Continuous example: Normal PDF is

$$f(y|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{y-\mu}{\sigma}\right)^2\right)$$
 - Discrete example: Poisson PMF is $f(y|\lambda) = \frac{\lambda^y \exp(-\lambda)}{y!}$
- The axioms of probability imply that

$$\underbrace{f(\theta|y)}_{\text{posterior}} = \frac{\overbrace{f(\theta)}^{\text{prior}} \overbrace{f(y|\theta)}^{\text{likelihood}}}{\underbrace{f(y)}_{\text{evidence}}}$$

Computation

- We typically have unnormalized posterior density
$$f(\theta|y) \propto f(\theta) f(y|\theta)$$
- Rather than getting analytical solution, the aim of inference is to draw samples from this unnormalized posterior density using Markov Chain Monte Carlo (MCMC).
- Draws are identically distributed but **not independent**, thus the name Markov Chain. As the sample size grows to infinity, the Markov Chain converges to the target distribution $f(\theta|y)$

Two General MCMC Engines

- Metropolis-Hastings algorithm. Initialize θ and repeat:
 - ① Randomly draw a new parameter vector, θ^* , from some jumping distribution $q(\theta^* | \theta, y)$
 - ② Evaluate $\frac{f(\theta^*|y)}{f(\theta|y)} \times \frac{q(\theta|\theta^*, y)}{q(\theta^*|\theta, y)} = \frac{f(\theta^*)f(y|\theta^*)}{f(\theta)f(y|\theta)} \times \frac{q(\theta|\theta^*, y)}{q(\theta^*|\theta, y)}$
 - ① If greater than a random draw from a standard uniform distribution, set $\theta = \theta^*$
 - ② Otherwise, retain $\theta = \theta$
 - ③ Optionally store θ after some warmup period
- Gibbs sampler
 - Can be seen as a special case of M-H
 - Update one parameter at a time; cycle through parameters
 - Specify $q(\theta_i^* | \theta_{-i}, y)$ as full-conditional PDF of i th parameter, given all other parameters θ_{-i}
 - Critical ratio always equals 1 so always accept proposals

Weaknesses of Metropolis-Hastings and Gibbs

- For M-H, a lot hinges on choice of $q(\theta^* | \theta, y)$
 - For “easy” jumping distributions, sampler randomly walks
 - Difficult to get to stationary distribution
 - Difficult to get into and out of tails of distribution
 - Tradeoff between acceptance probability and long jumps
 - High acceptance prob. \implies short jumps & high dependence
 - Long jumps \implies low acceptance prob. & high dependence
- For Gibbs, the problems are a bit different
 - $q(\theta_i^* | \theta_{-i}, y)$ can be hard to derive analytically
 - Conditional variance may be much smaller than marginal variance, which implies $\theta_i^* \approx \theta_i$
 - Thus, consecutive draws can have high dependence

Hamiltonian Monte Carlo (HMC) and Stan

- Very clever symmetric $q(\theta^* | \theta, y) = q(\theta | \theta^*, y)$
 - Based on metaphor of Hamiltonian dynamics
 - Solves Ordinary Differential Equations for θ^*
 - Allows long jumps with high acceptance probability
- Weaknesses of Hamiltonian Monte Carlo
 - Also need to compute $\nabla \theta$ (which is hard or slow)
 - Need to tune $q(\theta^* | \theta, y)$ during warmup period
- Strengths of Stan, which is a variant of HMC
 - Computes $\nabla \theta$ via automatic differentiation
 - Self-tuning, although you can also do it manually

Principles Needed to Use Stan Effectively

- In HMC, $\frac{q(\theta|\theta^*, y)}{q(\theta^*|\theta, y)} = 1$ so do not worry about that
- User needs to express in the Stan language
$$\ln(f(\theta) \times f(y|\theta)) = \ln(f(\theta)) + \ln(f(y|\theta))$$
- Can ignore terms that do not depend on θ
- You need to declare the support of θ
- HMC is vulnerable to varying parameter scales (some really big, others really small)
 - Stan mitigates this somewhat with tuning
 - User can mitigate it a lot with rescaling
- Stan is vulnerable to non-constant parameter dependence
 - Try to respecify your model in an equivalent way that reduces or regularizes the parameter dependence

How to use Stan

- 1 You write the model in (text) .stan file w/ R-like syntax
- 2 The parser, `stanc`, does two things checks that your model is valid and writes a conceptually equivalent C++ source file
- 3 C++ compiler (also linking `libstan`) creates a binary file from the C++ source
- 4 You execute the binary from R (or Python / command-line)
- 5 You summarize the resulting samples from the posterior
- 6 You conduct model checking and criticism

Types

- Primitive scalar types: `real` and `int`
- (column) `vector[K]` of `K` reals w/ 4 constrained subtypes
 - `simplex[K]` (non-negative and sums to 1)
 - `unit_vector[K]` (sum of squares equals 1)
 - `ordered[K]` (each element is greater than the previous)
 - `positive_ordered[K]` (and all are positive)
- `row_vector[K]` of reals
- `matrix[N,K]` of reals w/ 3 constrained subtypes
 - `cov_matrix[K]` (covariance matrix, or its inverse)
 - `cholesky_factor_cov[K,K]` (Cholesky factor thereof)
 - `corr_matrix[K]` (correlation matrix)
 - `cholesky_factor_corr[K]` (Cholesky factor thereof)
- `real`, `int`, `vector` (plain), `row_vector`, and `matrix` (plain) can have lower and / or upper bounds inside `<>`
- Can have homogenous arrays of any of the above, e.g.
`row_vector<lower=0,upper=1>[K] p[N];`

The data Block of a .stan File

- Contains everything passed from R to Stan
- Can be modeled data (\mathbf{y}), covariates (\mathbf{X}), constants (K)
- Basically, everything posterior distribution conditions on
- Can have comments in R style (#) or C++ style (// or /* */)

```
data {
  int<lower=1> K;    # number of covariates
  int<lower=1> N;    # number of observations
  matrix[N,K] X;    # predictor matrix
  real y[N];        # outcome variable

  // stuff for informative priors in regression
  vector[K] beta_0;
  cov_matrix[K] A_0;
  real<lower=0> v_0; /* or could be int */
  real<lower=0> s2_0;
}
```

Optional transformed data Block of a .stan File

- Is executed only once before the sampling iterations
- Can be used to calculate needed functions of data
- Not so necessary if calling Stan from R
- I often use it to check that data was passed correctly
- Need to declare objects before they are assigned (<-)

```
transformed data {  
  cov_matrix[K] XtXpA_0;  
  XtXpA_0 <- crossprod(X) + A_0;  
  print("K =", K);  
  print("N =", N);  
}
```

The parameters Block of a .stan File

- Declare everything whose posterior distribution is sought
- Cannot declare any integer parameters currently, only real
- Must specify the support of the parameters
- Stan is really sampling from unbounded parameter space
 - Behind-the-scenes transformations to yield parameters
 - For example, `exp()` of an unbounded yields a positive
 - Jacobians of **these** transformations handled automatically

```
parameters {  
  vector[K] beta;          # unrestricted  
  real<lower=0> sigma2; # restricted to be >= 0  
  /* legal to have lower and / or upper bounds  
  depend on the values of previously-declared  
  parameters */  
}
```

The optional transformed parameters Block

- Similar in structure to the transformed data block
- But is executed every iteration (and leapfrog step)
- Used to calculate deterministic functions of parameters
- Need to declare objects before they are assigned
- Such objects can then be used in the `model` block
- Constraints are validated and samples are stored

```
transformed parameters {  
  real<lower=0> sigma;  
  sigma <- sqrt(sigma2);  
}
```

The `model` Block of a `.stan` File

- Can declare more objects and then assign them
- Constraints are not validated and samples not stored
- Used to add log-priors and log-likelihood w/ `~` statements
- Can also manually increment the log-posterior

```
model {  
  y ~ normal(X * beta, sigma); # log-likelihood  
  beta ~ multi_normal_prec(beta_0, XtXpA_0);  
  sigma2 ~ inv_gamma(v_0, s2_0);  
}
```

The generated quantities Block of a .stan File

- Only evaluated for non-thinned post-warmup iterations
- Can declare more objects and then assign them
- Constraints are not validated but samples are stored
- Cannot reference anything in the `model` block
- Primarily used for
 - Interesting functions of posterior that don't go into likelihood
 - Posterior predictive distributions

```
generated quantities {  
  vector[N] y_tilde;  
  for (i in 1:N) {  
    y_tilde[i] <- normal_rng(row(X,i) * beta,  
                             sigma);  
  }  
}
```

Examples

- In RStan, any example can be executed with

```
posterior <- stan_demo() # choose example
```

- Also, can browse repo at

<https://github.com/stan-dev/stan/tree/develop/src/models> for same set of examples. Be sure to download

- foo.stan file
- foo.data.R file (invoke `pystan.misc.read_rdump("/path/to/foo.data.R")` to use with PyStan)

- Many examples in Stan manual at <http://mc-stan.org>