

W1211 Introduction to Statistics

Lecture 6

Wei Wang

Sep 24th, 2012

The Law of Total Probability

- The **Law of Total Probability** states, Let A_1, \dots, A_k be mutually exclusive and *exhaustive* events. Then for any other event B .

$$\begin{aligned} P(B) &= P(B|A_1)P(A_1) + \dots + P(B|A_k)P(A_k) \\ &= \sum P(B|A_i)P(A_i) \end{aligned}$$

- A_1, \dots, A_k are *exhaustive*, if one A_i must occur, so that $A_1 \cup \dots \cup A_k = S$.
- Proof: when $k=2$,

$$\begin{aligned} P(B) &= P((B \cap A) \cup (B \cap A^c)) \\ &= P(B \cap A) + P(B \cap A^c) \\ &= P(B|A)P(A) + P(B|A^c)P(A^c) \end{aligned}$$

Bayes Theorem

- With the help of the Law of Total Probability, we can state the Bayes Rule, which says, let A_1, \dots, A_k be a collection of k mutually exclusive and exhaustive events with *prior* probabilities $P(A_i)$ ($i=1, \dots, k$). Then for any other event B for which $P(B) > 0$, the *posterior* probability of A_j given that B has occurred is,

$$P(A_j|B) = \frac{P(A_j \cap B)}{P(B)} = \frac{P(B|A_j)P(A_j)}{\sum_{i=1}^k P(B|A_i)P(A_i)}$$

- When $k=2$, we have,

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|A^c)P(A^c)}$$

- Bayes Rule can be used to “*reverse*” the probability from the conditional probability that was originally given, or *to find the cause given the result*.

Bayes Theorem Example

- ▶ One percent of all individuals in a certain population are carriers of a particular disease. A diagnostic test for this disease has a 90% detection rate for carriers and a 5% detection rate for noncarriers. If a person is tested positive, what's the probability that this person is a carrier?

Bayes Theorem Example

- ▶ One percent of all individuals in a certain population are carriers of a particular disease. A diagnostic test for this disease has a 90% detection rate for carriers and a 5% detection rate for non-carriers. If a person is tested positive, what's the probability that this person is a carrier?

▶

$$\begin{aligned} &P(\text{is a carrier} | \text{tested positive}) \\ &= \frac{P(\text{carrier} \cap \text{tested positive})}{P(\text{tested positive})} \\ &= \frac{P(\text{positive} | \text{carrier})P(\text{carrier})}{P(\text{positive} | \text{carrier})P(\text{carrier}) + P(\text{positive} | \text{non-carrier})P(\text{non-carrier})} \end{aligned}$$

Independence

- ▶ Definition: Two events A and B are independent if $P(A|B) = P(A)$ (or alternatively $P(B|A) = P(B)$).

- ▶ A and B are independent if and only if

$$P(A \cap B) = P(A) \cdot P(B)$$

- ▶ Independent Events \neq Disjoint Events.

When will we have independence

- ▶ While, in the context of exam or homework problems, it is often given as the conditions.

When will we have independence

- ▶ While, in the context of exam or homework problems, it is often given as the conditions.
- ▶ Finite Population v.s. Infinite Population

Multiple Events

- Events A_1, \dots, A_n are **mutually independent** if for every k ($k = 2, 3, \dots, n$) and every subset of indices i_1, i_2, \dots, i_k ,

$$P(A_{i_1} \cap A_{i_2} \cap \dots \cap A_{i_k}) = P(A_{i_1}) P(A_{i_2}) \dots P(A_{i_k}).$$

- Independence is **very very important!**

Example

Ex. You recently bought a new set of tires from a manufacturer who just announced a recall because 2% of that particular brand were defective. What is the probability that at least one of your tires is defective? You may assume that the tires are defective independently of one another.

$$P(\text{at least one defective tire}) = 1 - P(\text{no defective tire})$$

Let A_i = tire i is not defective

$$P(A_i) = 1 - 0.02 = 0.98$$

$$\begin{aligned} P(\text{no defective tire}) &= P(A_1 \cap A_2 \cap A_3 \cap A_4) \\ &= P(A_1) P(A_2) P(A_3) P(A_4) = (0.98)^4 \end{aligned}$$

$$P(\text{at least one defective tire}) = 1 - (0.98)^4 = 0.0776$$

Random Variables

- ▶ A random variable is a variable whose value is a numerical outcome of a random phenomenon.
- ▶ For a given sample space S of some experiment, a random variable is any rule that associates a number with each outcome in S .
- ▶ To put it more mathematically, a random variable is a function whose domain is the sample space and whose range is the set of real numbers.

Random Variables v.s. Experiments

- ▶ An **experiment** is a physical setup in real world that provides us intuition about randomness.
- ▶ A **random variable** is a mathematical abstraction that describes randomness.
- ▶ When the outcome of the experiment can be seen as numerical, e.g., roll a die, we can effectively treat the experiment as a random variable.
- ▶ But for most RVs, especially continuous one, it is difficult to find some experiment that provides physical setup and intuition.

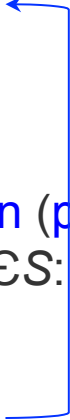
Discrete vs. Continuous

- X is a **discrete random variable** if its possible values either constitute a finite set or else can be listed in an infinite sequence in which there is a first element, a second element, and so on (“**countably**” infinite).
- X is a **continuous random variable** if it takes all possible values in an interval of numbers or all numbers in a disjoint union of such intervals. No possible value of the variable has positive probability, that is, $P(X=c) = 0$ for any possible value c .
- X can also be a random variable with a **mixture** distribution of both discrete and continuous components.

PMF

- The probability model for a discrete random variable X , lists its possible values and their probabilities.

Value of X	x_1	x_2	x_k
Probability	p_1	p_2	p_k

- Every probability, p_i , is a number between 0 and 1.
 - $p_1 + p_2 + \dots + p_k = 1$
 - The probability distribution or probability mass function (pmf) of a discrete rv is defined for every number x by $p(x) = P(X=x) = P(\text{all } s \in S: X(s)=x)$.
 - How to check if some function $p(x)$ is a proper PMF?
- 

Bernoulli RV

- ▶ The arguably simplest probability model is Bernoulli. Any random variable whose possible values are only 0 and 1 is called a Bernoulli random variable.
- ▶ Ex. Flip a coin. $S = \{H, T\}$. We can define a Bernoulli random variable, $X(H) = 1, X(T) = 0$. Then the distribution of X is

$$P(X = 1) = .5, P(X = 0) = .5$$

- ▶ Ex. Roll a die. $S = \{1, 2, 3, 4, 5, 6\}$. We can define a bernoulli random variable, $X(1) = X(2) = \mathbf{1}, X(3) = X(4) = X(5) = X(6) = \mathbf{0}$. Then the distribution is

$$P(X = \mathbf{1}) = 1/3, P(X = \mathbf{0}) = 2/3$$

Example

Ex. Flip three fair coins. (*Binomial*)

$S = \{\text{HHH, HHT, HTH, HTT, THT, THH, TTH, TTT}\}$. Let's define random variable X to be the number of heads in the experiment, i.e., $X(\text{HHH})=3$, $X(\text{THT})=1$, etc.

X

0 TTT

1 TTH THT HTT

2 THH HTH HHT

3 HHH

Value of X	0	1	2	3
Probability	0.125	0.375	0.375	0.125

One can calculate the probability of an event by adding the probabilities p_i of the particular values of x_i that make up the event. For example, if we want to know the probability of getting less than 2 heads, we can use

$$P(X < 2) = P(X=0) + P(X=1) = 0.125 + 0.375 = 0.5$$

$$\text{Note: } P(X \leq 2) = P(X=0) + P(X=1) + P(X=2) = 0.875$$

CDF


- The **cumulative distribution function** (cdf) $F(x)$ of a discrete rv variable X with pmf $p(x)$ is defined for every number x by

$$F(x) = P(X \leq x) = \sum_{y: y \leq x} p(y).$$

For any number x , $F(x)$ is the probability that the observed value of X will be at most x .

- For X a discrete rv, the graph of $F(x)$ will have a jump at every possible value of X and will be flat between possible values. Such a graph is called a **step function**.

The three coin flips example



Parameter and Family

- Suppose $p(x)$ depends on a quantity that can be assigned any one of a number of possible values, with each different value determining a different probability distribution. Such a quantity is called a **parameter** of the distribution. The collection of all probability distributions for different values of the parameter is called a **family** of probability distributions.

Ex. For Bernoulli rv's, the parameter is the probability of being 1 (or 0), that is,

$$p = P(X=1)$$

Expectation and Variance

- Random variables have distributions, so they have centers and spreads.
- The **expected value** (**mean value** or **expectation**) of a random variable describes its **theoretical long-run average value**.
- We typically use μ or $E(X)$ to denote the mean, $\text{Var}(X)$ to denote the variance and σ or $\text{SD}(X)$ to denote the standard deviation of a rv X .

Motivating examples

Ex. How many heads would you expect if you flipped a fair coin twice?

$S = \{\text{HH}, \text{HT}, \text{TH}, \text{TT}\}.$

$X =$ number of heads.

0 TT


1 HT TH

2 HH

$p(X=0) = 0.25; p(X=1) = 0.5; p(X=2) = 0.25.$

Each outcome is weighted by its probability.

$$\mu = 0 \times 0.25 + 1 \times 0.5 + 2 \times 0.25 = 1$$



Example

Ex. How many heads would you expect if you flipped a coin three times?

$$\mu = 0 \times 0.125 + 1 \times 0.375 + 2 \times 0.375 + 3 \times 0.125 = 1.5$$

This can never occur in a single trial of 3 flips. However, **on average** we would expect to get 1.5 heads if we repeated the experiment many times.

Definition

- Suppose X is a discrete random variable whose probability model is given by

Value of X	x_1	x_2	x_k
Probability	p_1	p_2	p_k

The expected value of X is given by

$$E(X) = \mu_X = \sum_{x \in D} x \cdot p(x) = x_1 p_1 + x_2 p_2 + \cdots x_k p_k$$

Example

Ex. Expectation of a Bernoulli rv.

$$p(x) = \begin{cases} 1-p & x=0 \\ p & x=1 \\ 0 & x \neq 0,1 \end{cases}$$

$$\mu = 0 \times (1-p) + 1 \times p = p.$$